

# Next generation smart contract & decentralised application platform, V. Buterin

Date 27/8/17 No.

- Satoshi solved distributed consensus by requiring proof of work from participants, as well as by implementing a public ledger
- Formally, a ledger can be thought of as a set of **state transitions**.

{  $APPLY(S, TX) \rightarrow S' \text{ or } ERROR$   
transaction converts old state into new state or throws error

- the state is the collection of **UTXO** (unspent transaction outputs), with each UTXO having a denomination and owner.
- a transaction maps 1+ inputs (with sigs) to 1+ outputs

for each input in TX:

- if input UTXO not in  $S$ :  $ERROR \leftarrow$  can't send coins that don't exist
- if sig  $\neq$  UTXO owner:  $ERROR \leftarrow$  can't send others' coins
- If  $\sum \text{inputs} < \sum \text{outputs}$ :  $ERROR \leftarrow$  can't create value
- return  $S - (\text{input UTXO}) + (\text{output UTXO})$

- Miners create blocks, which are valid if
  - previous block exists and is valid
  - timestamp  $>$  timestamp(previous)
  - valid proof of work
  - all transactions are valid
- An Eve with  $>51\%$  of the network's computational power could extend her own malicious forks.
- The blockchain history is stored in a **Merkle tree**
  - a node can compare a block header with a small part of the tree, which reduces required storage size
  - the **simplified payment verification (SPV)** protocol allows for 'light nodes' to verify POW on block headers and only download a small part of the tree.



- To build a consensus protocol, you can either build a new network or build it on top of the bitcoin blockchain
    - the former is difficult to implement and most applications will be too small
    - the latter is not scalable, as you cannot have 'light nodes'
  - Bitcoin does have a scripting language that can:
    - create 'safety deposit boxes' that requires an additional key to open
    - implement merchant escrow
    - support cross-cryptocurrency exchange
- But it has important limitations:
- not Turing-complete
  - lack of state: UTXOs are spent or unspent, which limits possibilities
  - Blockchain-blindness: a UTXO cannot see the nonce and prevHash, which could be good sources of randomness.

## Ethereum

- Blockchain with a built-in Turing complete programming language
- The state in Ethereum is made up of **accounts**, with each account having a 20B address and containing four fields:
  - nonce
  - **ether balance**
  - **contract code**
  - **storage** (empty by default)
- Accounts are either externally owned (controlled by private keys) else they are **contract accounts** (controlled by their contract code).

← ether is the digital currency used to pay transaction fees



- The Eth equivalent of a Btc transaction is a **message**:
  - can be created externally or by a contract
  - can explicitly contain data
  - if the recipient is a contract account, they can return a response → **messages can be used as functions**
- 'Transactions' in Eth refers to the signed data package that stores the message, the recipient, the signature, the quantity of eth, data, and **STARTGAS** and **GASPRICE**
  - to the miner
- to prevent infinite loops, you must pay a certain amount per computation. The limit of how much you will pay is **STARTGAS**. If a transaction runs out of gas, state changes revert except for gas fees. Spare gas is returned to the sender.
- Contracts in Eth are created with a different transaction format.
- Contracts are **first class citizens**, capable of doing anything that an external individual can.

$\text{APPLY}(S, \text{Tx}) \rightarrow S'$  works as follows:

1. Check if tx is well-formed with a valid sig
2. Subtract  $\text{STARTGAS} \times \text{GASPRICE}$  from sender and increment the sender's nonce
3. Initialise  $\text{GAS} = \text{STARTGAS}$ , and subtract a certain amount of gas per byte to pay for the size of the tx.
4. Transfer the tx value to the receiving account. If it is a contract, run the contract.
5. Miner collects the fees.

### Code execution

- The basis of Eth is **Ethereum Virtual Machine (EVM)** code
- Code is an infinite loop (incrementing a counter) until **STOP** or **RETURN** is seen.



- Each byte represents an operation. These operations can access:
    - the **stack** (32 B)
    - **memory**, an infinitely expandable byte array
    - the contract's **storage**, a key/value store where any item can be 32 B
  - value, sender, data, block-header
  - The code can also return a byte array
- reset after computation ends.
- data or more operations.

## The Ethereum Blockchain

- In addition to the transactions, blocks also contain:
    - the most recent state
    - the block number and difficulty.
- not inefficient because only a small part will change between transactions

### The block validation algorithm:

1. Check if the previous block exists and is valid
  2. Check the timestamp, block number, difficulty, tx root, gas limit
  3. Check the POW
  4. Set  $S[0] := \text{STATE\_ROOT}$  of the previous block.
  5.  $\text{APPLY}(S[i], \text{TX}[i])$  for  $i = 0, \dots, n-1$ .  $S\_FINAL := S[n]$
  6. Check if  $S\_FINAL$  is the same as the system's  $\text{STATE\_ROOT}$
- + miner reward

no need to store blockchain history

## Applications

- Token systems
- Financial derivatives, which require a data feed contract that can be pinged when needed.
- 'Decentralised dropbox': split the file and have your contract pay them as long as they can prove that they have the file.

- Decentralised Autonomous Organisations/corporations/communities (DAOs)
  - only a 67% majority of members can move funds / modify code.
  - there can be dividend-receiving shareholders and tradeable shares.
  - voting and liquid vote-delegation
- Decentralised data feeds :  $N$  parties input a datum, and everyone between the 25th-75th percentile gets a reward.
- Cloud computing : pay others to compute, with spotchecks built in.
- P2P gambling