

## Next generation smart contract & decentralised application platform, V. Buterin

Date 27/12/17 No.

- Satoshi solved distributed consensus by requiring proof of work from participants, as well as by implementing a public ledger
- Formally, a ledger can be thought of as a set of state transitions.

$\{ \text{APPLY}(S, TX) \rightarrow S' \text{ or ERROR}$

transaction converts old state into new state or throws error

- the state is the collection of UTXO (unspent transaction outputs), with each UTXO having a denomination and owner.
- a transaction maps 1+ inputs (with sigs) to 1+ outputs

for each input in TX:

if input UTXO not in  $S$ : ERROR  $\leftarrow$  can't send coins that don't exist

if sig  $\neq$  UTXO owner: ERROR  $\leftarrow$  can't send others' coins

If  $\sum \text{inputs} < \sum \text{outputs}$ : ERROR  $\leftarrow$  can't create value

return  $S - (\text{input UTXO}) + (\text{output UTXO})$

- Miners create blocks, which are valid if

- previous block exists and is valid
- timestamp  $>$  timestamp(previous)
- valid proof of work
- all transactions are valid

- An Eve with  $>51\%$  of the network's computational power could extend her own malicious forks.

- The blockchain history is stored in a Merkle tree

- a node can compare a block header with a small part of the tree, which reduces required storage size
- the simplified payment verification (SPV) protocol allows for 'light nodes' to verify POW on block headers and only download a small part of the tree.

- To build a consensus protocol, you can either build a new network or build it on top of the bitcoin blockchain
    - the former is difficult to implement and most applications will be too small
    - the latter is not scalable, as you cannot have 'light nodes'
  - Bitcoin does have a scripting language that can:
    - create 'safety deposit boxes' that requires an additional key to open
    - implement merchant escrow
    - support cross-cryptocurrency exchange
- But it has important limitations:
- not Turing-complete
  - lack of state: UTXOs are spent or unspent, which limits possibilities
  - Blockchain-blindness: a UTXO cannot see the nonce and prevHash, which could be good sources of randomness.

## Ethereum

- Blockchain with a built-in Turing complete programming language
  - The state in Ethereum is made up of accounts, with each account having a 20B address and containing four fields:
    - nonce
    - ether balance
    - contract code
    - storage (empty by default)
- ← ether is the digital currency used to pay transaction fees
- Accounts are either externally owned (controlled by private keys) else they are contract accounts (controlled by their contract code).

- The Eth equivalent of a Btc transaction is a message:
  - can be created externally or by a contract
  - can explicitly contain data
  - if the recipient is a contract account, they can return a response → messages can be used as functions
- 'Transactions' in Eth refers to the signed data package that stores the message, the recipient, the signature, the quantity of eth, data, and STARTGAS and GASPRICE → to the miner

to prevent infinite loops, you must pay a certain amount <sup>✓</sup> per computation. The limit of how much you will pay is STARTGAS. If a transaction runs out of gas, state changes revert except for gas fees. Spare gas is returned to the sender.

- Contracts in Eth are created with a different transaction format.
- Contracts are first class citizens, capable of doing anything that an external individual can.

$APPLY(S, Tx) \rightarrow S'$  works as follows:

1. Check if tx is well-formed with a valid sig
2. Subtract  $STARTGAS \times GASPRICE$  from sender and increment the sender's nonce
3. Initialise  $GAS = STARTGAS$ , and subtract a certain amount of gas per byte to pay for the size of the tx.
4. Transfer the tx value to the receiving account. If it is a contract, run the contract.
5. Miner collects the fees.

### Code execution

- The basis of Eth is Ethereum Virtual Machine (EVM) code
- Code is an infinite loop (incrementing a counter) until STOP or RETURN is seen.

- Each byte represents an operation. These operations can access:
    - the stack (32B)
    - memory, an infinitely expandable byte array
    - the contract's storage, a key/value store where any item can be 32B
    - value, sender, data, block-header
  - The code can also return a byte array
- } reset after computation ends.  
data or more operations.

## The Ethereum Blockchain

- In addition to the transactions, blocks also contain:
  - the most recent state ← not inefficient because only a small part will change between transactions
  - the block number and difficulty.

The block validation algorithm:

1. Check if the previous block exists and is valid
2. Check the timestamp, block number, difficulty, tx root, gas limit
3. Check the POW
4. Set  $S[0] := \text{STATE\_ROOT}$  of the previous block. + miner reward
5.  $\text{APPLY}(S[i], \text{TX}[i])$  for  $i = 0, \dots, n-1$ .  $S\_FINAL := S[n]$
6. Check if  $S\_FINAL$  is the same as the system's  $\text{STATE\_ROOT}$

no need to store blockchain history

## Applications

- Token systems
- Financial derivatives, which require a data feed contract that can be pinged when needed.
- 'Decentralised dropbox': split the file and have your contract pay them as long as they can prove that they have the file.

- Decentralised Autonomous Organisations/corporations/communities (DAOs)
  - only a 67% majority of members can move funds / modify code.
  - there can be dividend-receiving shareholders and tradeable shares.
  - voting and liquid vote-delegation
- Decentralised data feeds:  $N$  parties input a datum, and everyone between the 25th-75th percentile gets a reward.
- Cloud computing: pay others to compute, with spotchecks built in.
- P2P gambling